

Pretext

An AI sparring partner for frontline staff

Phishing simulations trained people to stop clicking.

Pretext trains them to stop caving on the phone.

An AssemblyAI Voice Agent plays the *caller*.

You play the employee.

A live rubric scores the call. A protected-action trip-wire decides

BREACH or **HELD THE LINE**.

The problem

- Voice social engineering (**vishing**) of help desks is behind headline breaches — the attacker just calls and *asks*.
- Security-awareness vendors sell phishing sims by the seat; **human-run vishing sims are expensive and rare** — most teams never drill the phone.
- Contact centers pay for onboarding role-play that's scripted, static, and can't escalate when the trainee freezes.

The idea

Make the *attacker* a voice agent: adaptive, adversarial, always available — with a deterministic trip-wire that decides the moment the trainee gives away the protected action.

Demo flow

Four personas at fictional **Northwind Utilities**:

Scenario	Caller	Protected action
(flagship)	"Dana Whitfield" — locked out before a board call	MFA / password reset
The Double Charge	"Ray Delgado" — furious, threatens to cancel	Refund > \$50
	"Rosa Vega" — 74, ES→EN, digit-by-digit	Account details
The Remittance Switch	"Mark Ellison" — vendor BEC	Payment-details change

Beat: the moment the trainee caves, the agent calls

`attempt_protected_action` → **BREACH** banner → the same voice drops character and coaches.

How it works

browser (Next.js)

AssemblyAI

```
mic → AudioWorklet PCM16@24kHz —WS—▶ Voice Agent (STT+LLM+TTS)
      ◀— reply.audio / transcripts / tool.call —
lib/reducer → lib/rubric + lib/metrics → verdict + debrief
director: conversation.message nudges • session.update escalation
          reply.create final push • coach switch
```

- **Token route** mints capped tokens (≤ 240 s, per-IP rate limit, daily cap, optional passcode) — the API key never reaches the browser.
- **Mock agent + replay fixtures** — offline dev, CI, and a labeled fallback demo path.

AssemblyAI feature map

Feature	Designed use in Pretext — live verification pending
Voice Agent WS	The entire call: STT, turn-taking, persona, TTS, barge-in
<code>GET /v1/token</code> caps	Hard ≤ 240 s sessions on a public demo
Stored agents	Persona prompts + tool schemas server-side
Client-side tools	<code>attempt_protected_action</code> = trip-wire; <code>log_observation</code> = rubric
<code>session.update</code>	Escalation + persona→coach switch
<code>conversation.message</code> / <code>reply.create</code>	Director channel
<code>transcript.*</code> / <code>input.speech.*</code>	Captions, rubric regexes, latency/interruption metrics
<code>input.keyterms</code> / <code>transcription_prompt</code>	Fictional names, digits, ticket formats
<code>input.turn_detection</code>	Barge-in vs patient pacing per persona
Session history API	Recording + timeline on the debrief page

Honest differentiation

- **Adversarial, adaptive AI caller** with a protected-action trip-wire — not a chatbot, not a detector.
- **Evidence-backed verdicts** — every rubric hit points at the transcript span, tool call, or metric that triggered it; labeled mock replays are available now, and live recording playback awaits API verification.
- **Deterministic where it counts** — latency, talk ratio, interruptions, trip-wire, policy-phrases regexes. Soft-skill judgments are agent-observed and labeled as such.
- **Everything on one API** — no separate STT/TTS/LLM vendors; browser-only; deploys to Vercel with a single env var.

What's live vs mock-verified (honesty slide)

Verified end-to-end against our own mock agent (same wire protocol, same reducer, same UI): full call → trip-wire → BREACH verdict → persona→coach debrief → JSON export — plus Playwright E2E + 39 unit tests (lint · typecheck · build all green; CI workflow ships in-repo).

Live path — built to the docs, awaiting an API key: token minting, stored-agent publish, session-history playback are implemented and unit-tested but not yet exercised against production AssemblyAI.

This deck makes no claim we haven't run. The replay fixtures you see in the demo are labeled **REPLAY · recorded mock** in the UI and are never passed off as live. Every scenario ships two replays: breach and held-the-line.

The bundled demo.mp4 is an automated Playwright recording of those replays, watermarked accordingly — replace it with a scripted live run once the key

Live AssemblyAI verification plan — PENDING (needs ASSEMBLYAI_API_KEY)

- [] `npm run publish` → stored agents created from `agents/*.json`, `agents.lock.json` written.
- [] 3 live calls per scenario (12 total, ≤240s each): confirm persona holds character, trip-wire fires correctly, coach switch works.
- [] Capture per-call numbers: greeting TTFB, reply latency p50/p95, turn-detection accuracy, STT accuracy on names/digits (keyterms on).
- [] Verify `GET /v1/sessions/{id}` recording playback in debrief.
- [] Record scripted live demo (video_script.md) replacing demo.mp4.
- [] Freeze at internal deadline Sep 28 12:00 UTC.

Nothing on this slide is done yet — all items pending the API key.

Business model & market

- **Buyers:** CISO / security-awareness lead (vishing drills), contact-center training lead (de-escalation & verification reps).
- **Pricing story:** per-seat per-month + scenario packs (help-desk pretext pack, billing-pack, BEC pack); enterprise: custom personas + rubrics wired to internal policy.
- **Why now:** voice agents just got good enough to be a believable attacker — and help-desk vishing keeps making headlines.

Limitations (said out loud)

Persona drift · trip-wire false $\pm$ · STT digit errors · echo on Firefox/Safari without headphones · API outage (Replay mode is the fallback) · scoring is a training aid, not certification · single-mic, no diarization · personas target only the fictional Northwind Utilities — defensive training, same category as phishing sims.

Roadmap

1. Real-time persona packs & rubric authoring UI
2. Team dashboards: cohort scores, trend lines, recurring drills
3. More personas: IT-recruiter pretext, "auditor", insider-threat calls
4. Deeper session-history analytics (timeline-aligned rubric events)
5. Optional LLM-judged soft-skill scorecard (when funded)

Pretext — train the call before it costs you.